


ISLR-HW5

Xiang XU

2/28/2019

```
##3
```

```
x = -2:2  
y = 1+x-2*(x-1)^2*I(x>1)  
plot(x,y)
```



```
##9
```

```
library(MASS)  
data(Boston)  
  
#a)  
reg = lm(nox~poly(dis,3),data=Boston)  
summary(reg)
```


```
##  
## Call:  
## lm(formula = nox ~ poly(dis, 3), data = Boston)  
##  
## Residuals:  
##      Min       1Q   Median       3Q      Max
```

```
## -0.121130 -0.040619 -0.009738 0.023385 0.194904
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.554695   0.002759 201.021 < 2e-16 ***
## poly(dis, 3)1 -2.003096   0.062071 -32.271 < 2e-16 ***
## poly(dis, 3)2  0.856330   0.062071  13.796 < 2e-16 ***
## poly(dis, 3)3 -0.318049   0.062071  -5.124 4.27e-07 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06207 on 502 degrees of freedom
## Multiple R-squared:  0.7148, Adjusted R-squared:  0.7131
## F-statistic: 419.3 on 3 and 502 DF,  p-value: < 2.2e-16
```

```
coef(reg)
```

```
## (Intercept) poly(dis, 3)1 poly(dis, 3)2 poly(dis, 3)3
## 0.5546951 -2.0030959 0.8563300 -0.3180490
```

```
reg.pred <- predict(reg,dis=list(Boston$dis))
plot(nox~dis,data=Boston)
lines(Boston$dis,reg.pred,col="red",lwd=2)
```



```
#b)
rss <- rep(NA,10)
for (i in 1:10){
  reg <- lm(nox~poly(dis,i),data=Boston)
```

```


    rss[i] <- sum(reg$residuals^2)
  }
  rss

## [1] 2.768563 2.035262 1.934107 1.932981 1.915290 1.878257 1.849484
## [8] 1.835630 1.833331 1.832171

# We can see that as i increases, the RSS decreases.

#c
library(boot)
error <- rep(NA,10)
for (i in 1:10){
  reg <- glm(nox~poly(dis,i),data=Boston)
  error[i] <- cv.glm(Boston,reg,K=10)$delta[2]
}
plot(1:10,error,xlab="Degree",ylab="CV error",type="line")

```



```

#From the plot, we may want to select degree of three.

#d
library(splines)
reg <- lm(nox~bs(dis,df=4,knots=c(4,7,11)),data=Boston)
summary(reg)

```


```

##
## Call:

```

```
## lm(formula = nox ~ bs(dis, df = 4, knots = c(4, 7, 11)), data = Boston)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.124567 -0.040355 -0.008702  0.024740  0.192920
##
## Coefficients:
##                                Estimate Std. Error t value Pr(>|t|)
## (Intercept)                   0.73926    0.01331  55.537 < 2e-16
## bs(dis, df = 4, knots = c(4, 7, 11))1 -0.08861    0.02504  -3.539  0.00044
## bs(dis, df = 4, knots = c(4, 7, 11))2 -0.31341    0.01680 -18.658 < 2e-16
## bs(dis, df = 4, knots = c(4, 7, 11))3 -0.26618    0.03147  -8.459 3.00e-16
## bs(dis, df = 4, knots = c(4, 7, 11))4 -0.39802    0.04647  -8.565 < 2e-16
## bs(dis, df = 4, knots = c(4, 7, 11))5 -0.25681    0.09001  -2.853  0.00451
## bs(dis, df = 4, knots = c(4, 7, 11))6 -0.32926    0.06327  -5.204 2.85e-07
##
## (Intercept)                    ***
## bs(dis, df = 4, knots = c(4, 7, 11))1 ***
## bs(dis, df = 4, knots = c(4, 7, 11))2 ***
## bs(dis, df = 4, knots = c(4, 7, 11))3 ***
## bs(dis, df = 4, knots = c(4, 7, 11))4 ***
## bs(dis, df = 4, knots = c(4, 7, 11))5 **
## bs(dis, df = 4, knots = c(4, 7, 11))6 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.06185 on 499 degrees of freedom
## Multiple R-squared:  0.7185, Adjusted R-squared:  0.7151
## F-statistic: 212.3 on 6 and 499 DF, p-value: < 2.2e-16


pred <- predict(reg,dis=list(Boston$dis))
plot(nox~dis,data=Boston)
lines(Boston$dis,pred,col="red")
```



```
#e
rss <- rep(NA,17)
for (i in 3:20){ #df should be greater than three
  reg <- lm(nox~bs(dis,df=i),data=Boston)
  rss[i] <- sum(reg$residuals^2)
}
rss[3:20]

## [1] 1.934107 1.922775 1.840173 1.833966 1.829884 1.816995 1.825653
## [8] 1.792535 1.796992 1.788999 1.782350 1.781838 1.782798 1.783546
## [15] 1.779789 1.775838 1.774487 1.776727

#f
error <- rep(NA,20)
for (i in 3:20){
  reg <- glm(nox~bs(dis,df=i),data=Boston)
  error[i] <- cv.glm(Boston,reg,K=10)$delta[2]
}
plot(3:20,error[3:20],xlab="df",ylab="CV error",type="l")
```



```
##10
```

```
#a
set.seed(1)
library(ISLR)
```

```
## Warning: package 'ISLR' was built under R version 3.5.2
```

```
library(leaps)
```


```
## Warning: package 'leaps' was built under R version 3.5.2
```

```
data("College")
train <- sample(length(College$Outstate),length(College$Outstate)/2)
Col.train <- College[train,]
Col.test <- College[-train,]
reg <- regsubsets(Outstate~.,data=Col.train,nvmax=17,method = "forward")
reg.sum <- summary(reg)
par(mfrow=c(1,3))
plot(reg.sum$cp,xlab="Number of Variables",ylab="Cp",type="l")
min.cp = min(reg.sum$cp)
std.cp = sd(reg.sum$cp)
abline(h = min.cp + 0.2 * std.cp, col = "red", lty = 2)
abline(h = min.cp - 0.2 * std.cp, col = "red", lty = 2)
plot(reg.sum$bic, xlab = "Number of Variables", ylab = "BIC", type = "l")
min.bic = min(reg.sum$bic)
std.bic = sd(reg.sum$bic)
abline(h = min.bic + 0.2 * std.bic, col = "red", lty = 2)
```

```

abline(h = min.bic - 0.2 * std.bic, col = "red", lty = 2)
plot(reg.sum$adjr2, xlab = "Number of Variables", ylab = "Adjusted R2",
     type = "l", ylim = c(0.4, 0.84))
max.adj2 = max(reg.sum$adjr2)
std.adj2 = sd(reg.sum$adjr2)
abline(h = max.adj2 + 0.2 * std.adj2, col = "red", lty = 2)
abline(h = max.adj2 - 0.2 * std.adj2, col = "red", lty = 2)

```



#From these plot, we may want to select 6 as the best subset size.

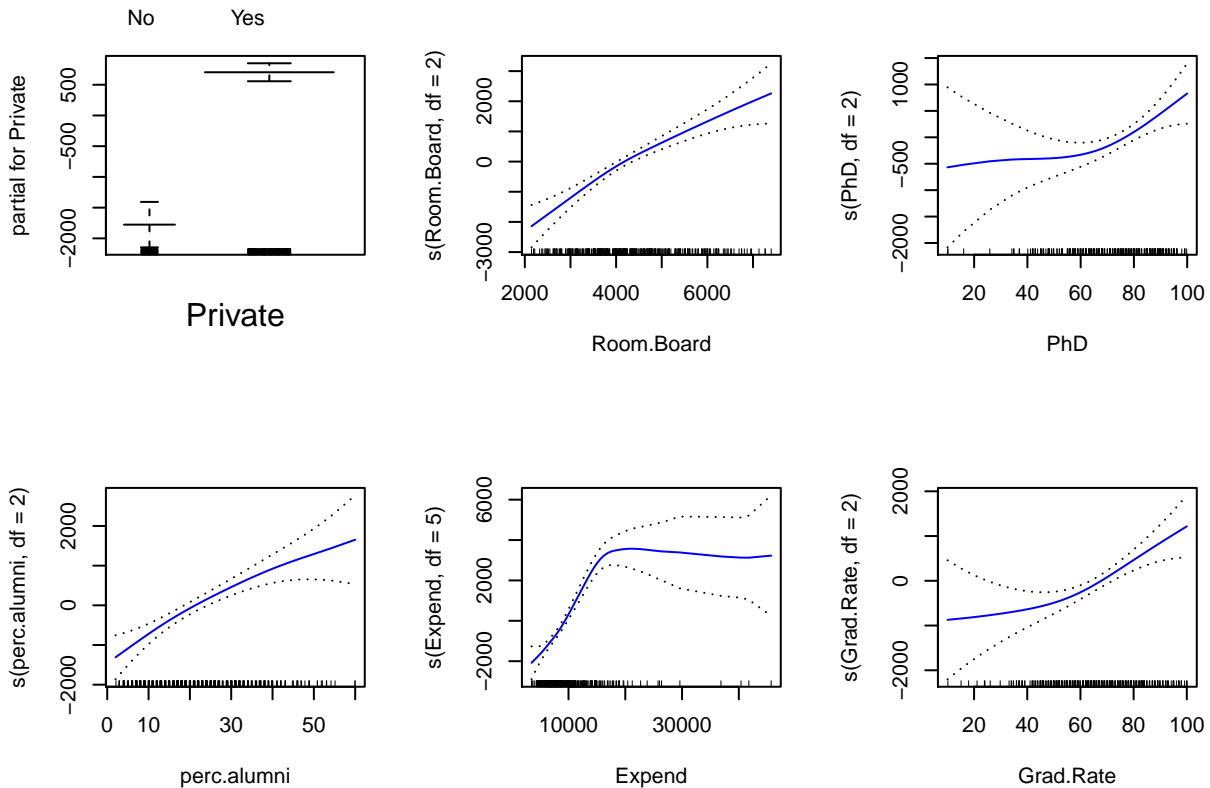
```

#b
library(gam)

## Warning: package 'gam' was built under R version 3.5.2
## Loading required package: foreach
## Loaded gam 1.16

reg <- gam(Outstate~Private+s(Room.Board,df=2)+s(PhD,df=2)+
          s(perc.alumni, df = 2) + s(Expend, df = 5) +
          s(Grad.Rate, df = 2), data = Col.train)
par(mfrow=c(2,3))
plot(reg,se=T,col="blue")

```



```
#c
pred <- predict(reg, Col.test)
error <- mean((Col.test$Outstate - pred)^2)
error

## [1] 3745460

tss = mean((Col.test$Outstate - mean(Col.test$Outstate))^2)
test.rss = 1 - error/tss
test.rss
```

```
## [1] 0.7696916
```

#Using 6 variables, we got a R-square of 77%.

```
#d
summary(reg)
```

```
##
## Call: gam(formula = Outstate ~ Private + s(Room.Board, df = 2) + s(PhD,
##      df = 2) + s(perc.alumni, df = 2) + s(Expend, df = 5) + s(Grad.Rate,
##      df = 2), data = Col.train)
## Deviance Residuals:
##      Min       1Q   Median       3Q      Max
## -4977.74 -1184.52   58.33  1220.04  7688.30
##
## (Dispersion Parameter for gaussian family taken to be 3300711)
##
```



```
## Null Deviance: 6221998532 on 387 degrees of freedom
## Residual Deviance: 1231165118 on 373 degrees of freedom
## AIC: 6941.542
##
## Number of Local Scoring Iterations: 2
##
## Anova for Parametric Effects
##
```

	Df	Sum Sq	Mean Sq	F value	Pr(>F)
Private	1	1779433688	1779433688	539.106	< 2.2e-16 ***
s(Room.Board, df = 2)	1	1221825562	1221825562	370.171	< 2.2e-16 ***
s(PhD, df = 2)	1	382472137	382472137	115.876	< 2.2e-16 ***
s(perc.alumni, df = 2)	1	328493313	328493313	99.522	< 2.2e-16 ***
s(Expend, df = 5)	1	416585875	416585875	126.211	< 2.2e-16 ***
s(Grad.Rate, df = 2)	1	55284580	55284580	16.749	5.232e-05 ***
Residuals	373	1231165118	3300711		

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Anova for Nonparametric Effects
##
```

	Npar	Df	Npar F	Pr(F)
(Intercept)				
Private				
s(Room.Board, df = 2)	1	3.5562	0.06010	.
s(PhD, df = 2)	1	4.3421	0.03786	*
s(perc.alumni, df = 2)	1	1.9158	0.16715	
s(Expend, df = 5)	4	16.8636	1.016e-12	***
s(Grad.Rate, df = 2)	1	3.7208	0.05450	.

```
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

*#From the Nonparametric Effects' Anova:
#A strong evidence of non-linear relationship between response variable and expend.*

```
#11
```

```
#a
set.seed(1)
X1 = rnorm(100)
X2 = rnorm(100)
eps = rnorm(100, sd = 0.1)
Y = -2.1 + 1.3 * X1 + 0.54 * X2 + eps


#b
beta0 = rep(NA, 1000)
beta1 = rep(NA, 1000)
beta2 = rep(NA, 1000)
beta1[1] = 18

#c
for (i in 1:1000) {
  a = Y - beta1[i] * X1
  beta2[i] = lm(a ~ X2)$coef[2]
  a = Y - beta2[i] * X2
  lm.fit = lm(a ~ X1)
  if (i < 1000) {
```

```

    beta1[i + 1] = lm.fit$coef[2]
  }
  beta0[i] = lm.fit$coef[1]
}
plot(1:1000, beta0, type = "l", xlab = "iteration", ylab = "betas", ylim = c(-2.2,
  1.6), col = "green")
lines(1:1000, beta1, col = "red")
lines(1:1000, beta2, col = "blue")
legend("center", c("beta0", "beta1", "beta2"), lty = 1, col = c("green", "red",
  "blue"))

```




#Notice that the coefficients quickly attain stable points.

```

#f
lm.fit = lm(Y ~ X1 + X2)
plot(1:1000, beta0, type = "l", xlab = "iteration", ylab = "betas", ylim = c(-2.2,
  1.6), col = "green")
lines(1:1000, beta1, col = "red")
lines(1:1000, beta2, col = "blue")
abline(h = lm.fit$coef[1], lty = "dashed", lwd = 3, col = rgb(0, 0, 0, alpha = 0.4))
abline(h = lm.fit$coef[2], lty = "dashed", lwd = 3, col = rgb(0, 0, 0, alpha = 0.4))
abline(h = lm.fit$coef[3], lty = "dashed", lwd = 3, col = rgb(0, 0, 0, alpha = 0.4))
legend("center", c("beta0", "beta1", "beta2", "multiple regression"), lty = c(1,
  1, 1, 2), col = c("green", "red", "blue", "black"))

```



```
#g  
# We only need one iteration to obtain a good approximation.
```